

EPISTEMIC PLANNING: RECENT ADVANCEMENTS AND FUTURE DIRECTIONS

Alessandro Burigana

Free University of Bozen-Bolzano

November 11th, 2025


The 35th International Conference
on Automated Planning and Scheduling,
Melbourne, Victoria, Australia

Classical Planning

Example (Blocks World)

- An initial configuration of blocks are piled up in stacks is given;
- The agent can move one block (at a time) from the top of a stack to another;
- From an initial configuration, the agent must move the blocks to achieve a desired one.

Initial state:



Actions $\text{move}(b, x, y)$:

- $\text{Pre}(\text{move}(b, x, y)) = \text{On}(b, x) \wedge \text{Clear}(b) \wedge \text{Clear}(y)$
- $\text{Eff}(\text{move}(b, x, y)) = \{\text{On}(b, y), \text{Clear}(x), \neg \text{On}(b, x), \neg \text{Clear}(y)\} \triangleright \top$

Epistemic Planning

Epistemic planning: enrichment of classical planning with notions of knowledge and belief.

- **Epistemic states** represent what the agents know/believe about the world and others' perspective of the world.
- **Epistemic actions** can change both the world and the knowledge/belief of the agents.
- Agents have to reason about each others' (higher-order) knowledge/beliefs to reach a shared goal.
- We move from a propositional, single-agent, fully observable, deterministic setting to an modal, multi-agent, partially observable, non-deterministic one.

Semantics for Epistemic Planning

We can define two main families of semantics for epistemic planning:

1 Sentential approaches:

- **Epistemic states**: sets of formulas called knowledge (or belief) bases.
- **Epistemic actions**: typically allow to modify a state by adding/deleting epistemic formulas (akin to classical actions).

2 Dynamic Epistemic Logic:

- **Epistemic states**: pointed Kripke models, where a set of possible worlds represents different perspectives of agents about a situation.
- **Epistemic actions**: pointed event models, where a set of possible events represents different agents' view of some information change.

EPISTEMIC LOGIC

Syntax

Let P be a finite set of propositional atoms and $Ag = \{1, \dots, n\}$ a finite set of agents. The language $\mathcal{L}_{P, Ag}$ of Epistemic Logic is given by the BNF:

Definition (Language of Epistemic Logic)

$$\phi ::= p \mid \neg\phi \mid \phi \wedge \phi \mid \Box_i \phi,$$

- Operator \Box_i : depending on the context, describes what agent i knows or believes.
- Dual operator \Diamond_i ($\equiv \neg\Box_i\neg$): describes what agent i considers to be possible or compatible.

Semantics

An **epistemic state** represents both **factual** information and what agents **know/believe**.

Definition (Epistemic Model)


An **epistemic model** is a triple $M = (W, R, L)$, where:

- $W \neq \emptyset$ is a finite set of **possible worlds**;
- $R : Ag \rightarrow 2^{W \times W}$ assigns to each agent i an **accessibility relation** R_i ; and
- $L : W \rightarrow 2^P$ assigns to each world a **label**, being a finite set of atoms.

Definition (Epistemic State)

An **epistemic state** is a pair (M, w_d) , where $w_d \in W$ is the **designated world**.

Example




Semantics

Definition (Truth)

Let $s = (M, w_d)$, where $M = (W, R, L)$, be an **epistemic state** and let $w \in W$:

- | | | |
|-----------------------------------|-----|--|
| $(M, w) \models p$ | iff | $p \in L(w)$ |
| $(M, w) \models \neg\phi$ | iff | $(M, w) \not\models \phi$ |
| $(M, w) \models \phi \wedge \psi$ | iff | $(M, w) \models \phi$ and $(M, w) \models \psi$ |
| $(M, w) \models \Box_i \phi$ | iff | $\forall v$ if $wR_i v$ then $(M, v) \models \phi$ |

Example



To Know or to Believe?

How can **epistemic states** represent the **knowledge** and the **beliefs** of agents?

→ We model them via **axioms**.

	Axiom	Frame Property	Knowledge	Belief
K	$\Box_i(\phi \rightarrow \psi) \rightarrow (\Box_i\phi \rightarrow \Box_i\psi)$	-	✓	✓
T	$\Box_i\phi \rightarrow \phi$	Reflexivity	✓	
D	$\Box_i\phi \rightarrow \Diamond_i\phi$	Serilality	✓	✓
4	$\Box_i\phi \rightarrow \Box_i\Box_i\phi$	Transitivity	✓	✓
5	$\neg\Box_i\phi \rightarrow \Box_i\neg\Box_i\phi$	Euclideaness	✓	✓


An **epistemic state** represents:

- **Knowledge**, when it satisfies axioms **K**, **T**, **4** and **5** ⇒ **Logic S5_n**
- **Belief**, when it satisfies axioms **K**, **D**, **4** and **5** ⇒ **Logic KD45_n**

Epistemic Blocks World

Example (Epistemic Blocks World)


- Agent a : only sees from above.



Epistemic Blocks World

Example (Multi-Agent Epistemic Blocks World)


- Agent a : only sees from above.
- Agent l : only sees from a top left position.



Epistemic Blocks World

Example (Multi-Agent Epistemic Blocks World)


- Agent a : only sees from above.
- Agent r : only sees from a top right position.



Epistemic Blocks World

Example (Multi-Agent Epistemic Blocks World)

- Agent a : only sees from above.
- Agent l : only sees from a top left position.
- Agent r : only sees from a top right position.



DYNAMIC EPISTEMIC LOGIC

Epistemic Actions

Definition (Event Model)

An **event model** is a quadruple $A = (E, Q, \text{pre}, \text{post})$, where:

- $E \neq \emptyset$ is a finite set of **events**;
- $Q : \text{Ag} \rightarrow 2^{E \times E}$ assigns to each agent i an **accessibility relation** Q_i ;

Intuitively:

- An **event** can be seen as a **classical action**.
- Accessibility relations specify the perspectives of agents on which **events** take place.

Epistemic Actions

Definition (Event Model)

An **event model** is a quadruple $A = (E, Q, \text{pre}, \text{post})$, where:

- $E \neq \emptyset$ is a finite set of **events**;
- $Q : \text{Ag} \rightarrow 2^{E \times E}$ assigns to each agent i an **accessibility relation** Q_i ;
- $\text{pre} : E \rightarrow \mathcal{L}_{P, \text{Ag}}$ assigns to each event a **precondition**;
- $\text{post} : E \times P \rightarrow \mathcal{L}_{P, \text{Ag}}$ assigns to each event-atom pair a **postcondition**.

Intuitively:

- An **event** can be seen as a **classical action**, **each with its own pre- and postconditions**.
- Accessibility relations specify the perspectives of agents on which **events** take place.

Epistemic Actions

Definition (Event Model)

An **event model** is a quadruple $A = (E, Q, \text{pre}, \text{post})$, where:

- $E \neq \emptyset$ is a finite set of **events**;
- $Q : \text{Ag} \rightarrow 2^{E \times E}$ assigns to each agent i an **accessibility relation** Q_i ;
- $\text{pre} : E \rightarrow \mathcal{L}_{P, \text{Ag}}$ assigns to each event a **precondition**;
- $\text{post} : E \times P \rightarrow \mathcal{L}_{P, \text{Ag}}$ assigns to each event-atom pair a **postcondition**.

Intuitively:

- An **event** can be seen as a **classical action**, **each with its own pre- and postconditions**.
- Accessibility relations specify the perspectives of agents on which **events** take place.

Definition (Epistemic Action)

An **epistemic action** is a pair (A, e_d) , where $e_d \in E$ is the **designated event**.

Product Update

An **action** (A, e_d) is **applicable** in an **epistemic state** (M, w_d) iff $(M, w_d) \models pre(e_d)$.

Definition (Product Update)

Given (M, w_d) and (A, e_d) , where $M = (W, R, L)$ and $A = (E, Q, pre, post)$, their **product update** $(M, w_d) \otimes (A, e_d)$ is the **epistemic state** $((W', R', L'), w'_d)$ where:

Product Update

An **action** (A, e_d) is **applicable** in an **epistemic state** (M, w_d) iff $(M, w_d) \models \text{pre}(e_d)$.

Definition (Product Update)

Given (M, w_d) and (A, e_d) , where $M = (W, R, L)$ and $A = (E, Q, \text{pre}, \text{post})$, their **product update** $(M, w_d) \otimes (A, e_d)$ is the **epistemic state** $((W', R', L'), w'_d)$ where:

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;

Product Update

An **action** (A, e_d) is **applicable** in an **epistemic state** (M, w_d) iff $(M, w_d) \models \text{pre}(e_d)$.

Definition (Product Update)

Given (M, w_d) and (A, e_d) , where $M = (W, R, L)$ and $A = (E, Q, \text{pre}, \text{post})$, their **product update** $(M, w_d) \otimes (A, e_d)$ is the **epistemic state** $((W', R', L'), w'_d)$ where:

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_i = \{((w, e), (v, f)) \in W' \times W' \mid wR_i v \text{ and } eQ_i f\}$;

Product Update

An **action** (A, e_d) is **applicable** in an **epistemic state** (M, w_d) iff $(M, w_d) \models pre(e_d)$.

Definition (Product Update)

Given (M, w_d) and (A, e_d) , where $M = (W, R, L)$ and $A = (E, Q, pre, post)$, their **product update** $(M, w_d) \otimes (A, e_d)$ is the **epistemic state** $((W', R', L'), w'_d)$ where:

- $W' = \{(w, e) \in W \times E \mid (M, w) \models pre(e)\}$;
- $R'_i = \{((w, e), (v, f)) \in W' \times W' \mid wR_i v \text{ and } eQ_i f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models post(e, p)\}$; and

Product Update

An **action** (A, e_d) is **applicable** in an **epistemic state** (M, w_d) iff $(M, w_d) \models pre(e_d)$.

Definition (Product Update)

Given (M, w_d) and (A, e_d) , where $M = (W, R, L)$ and $A = (E, Q, pre, post)$, their **product update** $(M, w_d) \otimes (A, e_d)$ is the **epistemic state** $((W', R', L'), w'_d)$ where:

- $W' = \{(w, e) \in W \times E \mid (M, w) \models pre(e)\}$;
- $R'_i = \{((w, e), (v, f)) \in W' \times W' \mid wR_i v \text{ and } eQ_i f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models post(e, p)\}$; and
- $w'_d = (w_d, e_d)$.

Public Announcements

Example

Public Announcement

Agent r publicly tells everybody that
 $\neg On(b_1, s_3)$.

a, l, r


 $e : \langle \neg On(b_1, s_3), id \rangle$

Public Announcements

Example


Public Announcement

Agent r publicly tells everybody that
 $\neg \text{On}(b_1, s_3)$.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_r = \{((w, e), (v, f)) \in W' \times W' \mid w R_r v \text{ and } e Q_r f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.




Public Announcements

Example


Public Announcement

Agent r publicly tells everybody that
 $\neg \text{On}(b_1, s_3)$.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_r = \{((w, e), (v, f)) \in W' \times W' \mid wR_r v \text{ and } eQ_r f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.




Public Announcements

Example


Public Announcement

Agent r publicly tells everybody that
 $\neg \text{On}(b_1, s_3)$.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_r = \{((w, e), (v, f)) \in W' \times W' \mid wR_r v \text{ and } eQ_f f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.




Semi-Private Sensing Action

Example

Semi-Private Sensing Action

Agent **r** **peeks** under block b_2 while agents **a** and **l** observe him. Specifically:

- Agents **r** and **l** observe what is actually being sensed.
- Agent **a** can not directly observe what agent **r** is seeing.



Trivial postconditions are omitted.


Semi-Private Sensing Action

Example


Semi-Private Sensing Action

Agent **r** **peeks** under block b_2 while agents **a** and **l** observe him. Specifically:

- Agents **r** and **l** observe what is actually being sensed.
- Agent **a** can not directly observe what agent **r** is seeing.



Trivial postconditions are omitted.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_f = \{((w, e), (v, f)) \in W' \times W' \mid wR_f v \text{ and } eQ_f f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.


Semi-Private Sensing Action

Example



Semi-Private Sensing Action

Agent **r** **peeks** under block b_2 while agents **a** and **l** observe him. Specifically:

- Agents **r** and **l** observe what is actually being sensed.
- Agent **a** can not directly observe what agent **r** is seeing.



Trivial postconditions are omitted.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_f = \{((w, e), (v, f)) \in W' \times W' \mid wR_f v \text{ and } eQ_f f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.


Semi-Private Sensing Action

Example


Semi-Private Sensing Action

Agent r peeks under block b_2 while agents a and l observe him. Specifically:

- Agents r and l observe what is actually being sensed.
- Agent a can not directly observe what agent r is seeing.



Trivial postconditions are omitted.



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_d = \{((w, e), (v, f)) \in W' \times W' \mid w R_d v \text{ and } e Q_d f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.

Private Ontic Actions

Example

Private Ontic Action

Agent I privately moves block b_2 from b_1 to b_3 , where:

- $pre = On(b_2, b_1) \wedge Clear(b_2) \wedge Clear(b_3)$
- $post(e_1, On(b_2, b_1)) = \perp$
- $post(e_1, On(b_2, b_3)) = \top$




Private Ontic Actions

Example

Private Ontic Action

Agent $|$ privately moves block b_2 from b_1 to b_3 , where:

- $\text{pre} = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_3)$
- $\text{post}(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}(e_1, \text{On}(b_2, b_3)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_I = \{((w, e), (v, f)) \in W' \times W' \mid w R_I v \text{ and } e Q_I f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$


Private Ontic Actions

Example

Private Ontic Action

Agent $|$ privately moves block b_2 from b_1 to b_3 , where:

- $\text{pre} = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_3)$
- $\text{post}(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}(e_1, \text{On}(b_2, b_3)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\text{I} = \{((w, e), (v, f)) \in W' \times W' \mid w R_\text{I} v \text{ and } e Q_\text{I} f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$


Private Ontic Actions

Example

Private Ontic Action

Agent $|$ privately moves block b_2 from b_1 to b_3 , where:

- $\text{pre} = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_3)$
- $\text{post}(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}(e_1, \text{On}(b_2, b_3)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\}$;
- $R'_I = \{((w, e), (v, f)) \in W' \times W' \mid w R_I v \text{ and } e Q_I f\}$;
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\}$; and
- $w'_d = (w_d, e_d)$.


Private Ontic Actions

Example

Private Ontic Action

Agent $|$ privately moves block b_2 from b_1 to b_3 , where:

- $\text{pre} = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_3)$
- $\text{post}(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}(e_1, \text{On}(b_2, b_3)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\perp = \{((w, e), (v, f)) \in W' \times W' \mid w R_\perp v \text{ and } e Q_\perp f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$



Private Ontic Actions

Example

Private Ontic Action

Agent $|$ privately moves block b_2 from b_1 to b_3 , where:

- $\text{pre} = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_3)$
- $\text{post}(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}(e_1, \text{On}(b_2, b_3)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\perp = \{((w, e), (v, f)) \in W' \times W' \mid w R_\perp v \text{ and } e Q_\perp f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$

A Glance at Non-Deterministic Actions

Example

Public Non-Deterministic Ontic Action

Agents **publicly flip** a coin: if heads, they move b_2 from b_1 to b_3 , otherwise to b_4 , where:

- $\text{pre}_k = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_k)$
- $\text{post}_k(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}_k(e_1, \text{On}(b_2, b_k)) = \top$



$e_1 : \langle \text{pre}_3, \text{post}_3 \rangle$ $e_2 : \langle \text{pre}_4, \text{post}_4 \rangle$

A Glance at Non-Deterministic Actions


Example

Public Non-Deterministic Ontic Action

Agents **publicly flip** a coin: if heads, they move b_2 from b_1 to b_3 , otherwise to b_4 , where:

- $\text{pre}_k = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_k)$
- $\text{post}_k(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}_k(e_1, \text{On}(b_2, b_k)) = \top$

$$\begin{array}{cc} a, l, r & a, l, r \\ \text{e}_1 : \langle \text{pre}_3, \text{post}_3 \rangle & \text{e}_2 : \langle \text{pre}_4, \text{post}_4 \rangle \end{array}$$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\perp = \{((w, e), (v, f)) \in W' \times W' \mid wR_\perp v \text{ and } eQ_\perp f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$



A Glance at Non-Deterministic Actions

Example


Public Non-Deterministic Ontic Action

Agents **publicly flip** a coin: if heads, they move b_2 from b_1 to b_3 , otherwise to b_4 , where:

- $\text{pre}_k = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_k)$
- $\text{post}_k(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}_k(e_1, \text{On}(b_2, b_k)) = \top$



(v_1, e_1)



(v_1, e_2)

Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\text{f} = \{((w, e), (v, f)) \in W' \times W' \mid wR_\text{f} v \text{ and } eQ_\text{f} f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$



A Glance at Non-Deterministic Actions

Example


Public Non-Deterministic Ontic Action

Agents **publicly flip** a coin: if heads, they move b_2 from b_1 to b_3 , otherwise to b_4 , where:

- $\text{pre}_k = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_k)$
- $\text{post}_k(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}_k(e_1, \text{On}(b_2, b_k)) = \top$



(v_1, e_1)



(v_1, e_2)

Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_f = \{((w, e), (v, f)) \in W' \times W' \mid wR_f v \text{ and } eQ_f f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$


A Glance at Non-Deterministic Actions

Example

Public Non-Deterministic Ontic Action

Agents **publicly flip** a coin: if heads, they move b_2 from b_1 to b_3 , otherwise to b_4 , where:

- $\text{pre}_k = \text{On}(b_2, b_1) \wedge \text{Clear}(b_2) \wedge \text{Clear}(b_k)$
- $\text{post}_k(e_1, \text{On}(b_2, b_1)) = \perp$
- $\text{post}_k(e_1, \text{On}(b_2, b_k)) = \top$



Definition (Product Update)

- $W' = \{(w, e) \in W \times E \mid (M, w) \models \text{pre}(e)\};$
- $R'_\perp = \{((w, e), (v, f)) \in W' \times W' \mid wR_\perp v \text{ and } eQ_\perp f\};$
- $L'((w, e)) = \{p \in P \mid (M, w) \models \text{post}(e, p)\};$ and
- $w'_d = (w_d, e_d).$

The Epistemic Plan Existence Problem


(Epistemic)
Planning Task

	Initial Epistemic State	Set of Epistemic Actions	Goal Formula
--	-----------------------------------	------------------------------------	-------------------------


The Epistemic Plan Existence Problem



The Epistemic Plan Existence Problem



The Epistemic Plan Existence Problem



Epistemic Plan Existence Problem

Given an epistemic planning task, **does there exist a plan that achieves the goal?**

Classical Vs. Epistemic Actions

Classical planning:

- 1 Propositional
- 2 Single-agent
- 3 Fully Observable
- 4 Deterministic
- 5 Ontic change

Epistemic planning:

- 1 Modal
- 2 Multi-agent
- 3 Partially Observable
- 4 Non-deterministic (multi-pointed models are needed)
- 5 Ontic and epistemic change

Moreover, agents can reason on **higher-order knowledge/beliefs** of others to any nesting level.

→ There are **no bounds** on the reasoning power of agents!

Theorem (Bolander and Andersen [BA11])

The epistemic plan existence problem is undecidable.

RECENT ADVANCEMENTS

Several Different Directions

Many approaches have been pursued in epistemic planning. Today we cover the following:

- 1 Sentential approaches:
 - Compilations to classical planning.
 - Alternating cover disjunctive formulas.
- 2 Heuristics for epistemic planning.
- 3 DEL-based approaches:
 - (Bounded) bisimulation contractions.
 - Depth-bounded epistemic planning.

SENTENTIAL APPROACHES

Compilations to Classical Planning

PDDL translation by Kominis and Geffner [KG15] of the next epistemic planning formalism:

- **Public ontic actions**: all agents know both about the action and its effects.
- **Semi-private sensing/announcements actions**: all agents know about the action, but only some know its effects.

They show that:

- The compilation is quadratic.
- The identified fragment is **PSPACE-complete**.
- Their formalism corresponds to a **DEL fragment**.

Compilations to Classical Planning (cont.)

PDDL translation by Muise et al. [Mui+15; Mui+22] based on a restricted language:

Definition (Restricted Modal Literals)

$$\phi ::= p \mid \neg\phi \mid \Box_i \phi$$

- **Epistemic states:** sets of RMLs.
- **Epistemic actions:** preconditions/effects pairs defined over RMLs.
- **Promising results in different epistemic planning benchmarks.**
- Worse performances on instances with higher reasoning depth.
- More expressive than Kominis and Geffner's approach (e.g., allows for private actions).

Compilations to Classical Planning (cont.)

A similar approach is pursued by Cooper et al. [Coo+16], later generalized by [Coo+20]:

Definition (Epistemic Logic of Observation (EL-O))

$$\begin{aligned}\alpha ::= & \ p \mid S_i \alpha \mid JS\alpha \\ \phi ::= & \ \alpha \mid \neg\phi \mid \phi \wedge \phi\end{aligned}$$

where $S_i\phi$ means that agent i **sees whether** ϕ holds and $JS\phi$ means that all agents **jointly see** whether ϕ .

- $\phi \wedge S_i\phi$ is equivalent to $\square_i\phi$.
- $\phi \wedge JS\phi$ is equivalent to $C\phi$ (common knowledge of ϕ).
- They show that the problem is **PSPACE-complete**.
- More expressive than Muise et al.'s approach (allows for common knowledge and parallel actions).

Pros and Cons

Pros

- 👍 Rely on **efficiency** of classical planners.
- 👍 **Lower complexity** of the plan existence problem.

Cons

- 👎 **Limited** to specific fragments.
- 👎 Typically **do not scale well** when higher-order knowledge is involved.

Alternating Cover Disjunctive Formulas

Huang et al. [Hua+17] proposed a **doxastic planning framework**, i.e., based on the logic of belief KD45_n :

- **Epistemic states** are general KD45_n formulas with common knowledge.
- **Deterministic actions**: precondition/effects pairs over KD45_n formulas.
- **Sensing actions**: precondition + positive and negative effects over KD45_n formulas.
- Formulas are transformed into equivalent **Alternating Cover Disjunctive Formulas** (ACDF).
 - Length of an ACDF formula is shown to be **at most singly exponential** in the length of the original formula.

Alternating Cover Disjunctive Formulas: Pros and Cons

- A **Pruning AND-OR** (PrAO) search algorithm with visited state check is provided.
- The algorithm builds an action tree branching on sensing actions.
- Here a stronger notion of equivalence of ACDF formulas is introduced, which can be checked in polynomial time.
- The planner, called **MEPK**, is compared to the solvers by Kominis and Geffner, and by Muise et al.

Pros

- ❑ The formalism is **more expressive** than the compilation-based ones.
- ❑ **Reasonable** performances on the conducted experiments.

Cons

- ❑ **Worse performances** compared to compilation-based approaches.
- ❑ **Exponential blowup** of ACDF formulas size.

HEURISTICS FOR EPISTEMIC PLANNING

Heuristics for MEPK

Later, the MEPK planner was improved as follows [Wu18]:

- A normal form for ACDF formulas is provided, called **ADNF**, which is claimed to be more space efficient than regular ACDF.
 - A notion of **distance** is provided for ADNF formulas, used to guide the search towards states with lower distance from the goal.
 - Two heuristic strategies for pruning the search space are also provided.
- The resulting planner, called **MEPL**, was benchmarked against MEPK, showing improvements on the vast majority of the tested instances.

More Heuristics for MEPK

Heuristics for MEPK have been also developed in a subsequent work [FL24]:

- **Enhancement:** use information in the path leading to the first goal-satisfying state to guide the rest of the search.
- **Belief lock:** in some cases, once an agent has acquired some belief, it can not later forget it (the belief is “locked”).
 - A set of conditions is identified for **recognizing locked beliefs**.
 - Belief locks are used for **pruning** the search space.
 - For instance, if in the current state $\Box_i p$ is recognized as a locked belief, and the goal requires that $\Box_i \neg p$ instead, we can safely prune the search, as the goal is unreachable.
- The comparison with the original MEPK planner showed performance improvements. However, no comparison with MEPL was conducted.

Heuristics for EFP 2.0

EFP 2.0 was later equipped with heuristic search strategies [Fab+24]

- **Several heuristics** were proposed, based on planning graph methods and on maximal goal sub-formulas satisfaction.
- A **portfolio-like technique** was used to construct a machine learning model for selecting the best heuristic for each input problem.

Preliminary experiments showed the following:

- 👍 General **improvements** over EFP 2.0 with no heuristics.
- 👎 Minimality of plans **not guaranteed**.

DEL-BASED APPROACHES

Main Challenges

- Higher uncertainty of agents means bigger models.
- Worst-case exponential blowup of size of states after product update.
- Expensive check for visited states.
- Search space can be infinite.

Bisimulations

- A **bisimulation** between two states s and s' is a binary relation Z on their world-sets s.t.:
 - If $(x, x') \in Z$, then x and x' are **propositionally equivalent** (**atom**) and **for each i -successor y of x there exists an i -successor y' of x' s.t. $(x', y') \in Z$** (**forth**), and vice versa (**back**).
- If such a Z exists, we say that s and s' are **bisimilar**, written $s \Leftrightarrow s'$.
- Bisimilarity corresponds to **modal equivalence**:

Proposition ([BRV01])

Two states are **bisimilar** iff they **satisfy the same formulas in $\mathcal{L}_{P,\text{Ag}}$** .

Bisimulations

- A **bisimulation** between two states s and s' is a binary relation Z on their world-sets s.t.:
 - If $(x, x') \in Z$, then x and x' are **propositionally equivalent** (**atom**) and **for each i -successor y of x there exists an i -successor y' of x' s.t. $(x', y') \in Z$ (**forth**), and vice versa (**back**).**
- If such a Z exists, we say that s and s' are **bisimilar**, written $s \Leftrightarrow s'$.
- Bisimilarity corresponds to **modal equivalence**:

Proposition ([BRV01])

Two states are **bisimilar** iff they **satisfy the same formulas** in $\mathcal{L}_{P,\text{Ag}}$.

Example (Two bisimilar states)



Bisimulations

- A **bisimulation** between two states s and s' is a binary relation Z on their world-sets s.t.:
 - If $(x, x') \in Z$, then x and x' are **propositionally equivalent** (**atom**) and **for each i -successor y of x there exists an i -successor y' of x' s.t. $(x', y') \in Z$ (**forth**), and vice versa (**back**).**
- If such a Z exists, we say that s and s' are **bisimilar**, written $s \doteq s'$.
- Bisimilarity corresponds to **modal equivalence**:

Proposition ([BRV01])

Two states are **bisimilar** iff they **satisfy the same formulas in $\mathcal{L}_{P,\text{Ag}}$** .

Proposition (Product Update Preserves Bisimilarity [DHK07])

If $s \doteq s'$ and α is applicable in both, then:

$$s \otimes \alpha \doteq s' \otimes \alpha$$

Reducing the Size of Visited States

Definition (Bisimulation Contraction)

The **(bisimulation) contraction** of s is the **quotient structure** $[s]_{\equiv}$ of s induced by the bisimilarity relation.

Proposition ([BRV01])

$[s]_{\equiv}$ is a **minimal state** (*smallest number of worlds and edges*) **bisimilar** to s .

Reducing the Size of Visited States

Definition (Bisimulation Contraction)

The **(bisimulation) contraction** of s is the **quotient structure** $[s]_{\equiv}$ of s induced by the bisimilarity relation.

Proposition ([BRV01])

$[s]_{\equiv}$ is a **minimal state** (smallest number of worlds and edges) **bisimilar** to s .

💡 Key Idea

We can **replace any visited state s with** its bisimulation contraction $[s]_{\equiv}$.

- $[s]_{\equiv}$ and s are bisimilar, and so are $[s]_{\equiv} \otimes \alpha$ and $s \otimes \alpha$.
- The **size of $[s]_{\equiv}$ is at most the size of s** .
- We compute and store less information.
- Technique adopted by several epistemic planners [Fab+20; BDH21; BBM25].

EFP and PG-EFP

One of the earlier DEL-based planners is the **Epistemic Forward Planner** [Le+18], based on the $m\mathcal{A}^*$ epistemic action description language [Bar+15; Bar+22]:

- **Private/public ontic actions.**
- **Public/(Semi-)private sensing/announcements actions.**

Two search strategies were implemented:

- **EFP:** Breadth-First Search.
- **PG-EFP:** **planning graph heuristic** tailored for the $m\mathcal{A}^*$ fragment.

Later, Fabiano et al. [Fab+20] implemented an improved version of the planner, called **EFP 2.0**.

Two new search algorithms

- Kripke-based BFS search with **bisimulation contractions** and check for visited states.
- BFS search based on an alternative semantics for epistemic states called **possibilities**.
 - More **compact representation** of states.
 - Natural **reuse of previously calculated information**.

Later, Fabiano et al. [Fab+20] implemented an improved version of the planner, called **EFP 2.0**.

Two new search algorithms

- Kripke-based BFS search with **bisimulation contractions** and check for visited states.
- BFS search based on an alternative semantics for epistemic states called **possibilities**.
 - More **compact representation** of states.
 - Natural **reuse of previously calculated information**.

Experiments results

- Improved performances wrt. EFP 1.0, especially the possibility-based planner.
- Promising results in many epistemic planning benchmarks.
- Muise et al. [Mui+22] compared to EFP 2.0
 - They showed **better performances** than EFP 2.0 on **smaller instances**.
 - And **worse** results on instances that required **higher reasoning depth**.

Pros and Cons

Pros

- 👍 **Efficient** running times for the considered fragment.
- 👍 **Good scalability** on bigger instances.

Cons

- 👎 **Limited** to specific fragments.
- 👎 **Check for visited states** is computationally **expensive**.

Improving Check for Visited States

How do we check if a state s has already been visited? Suppose we have a set Visited of visited states:

Improving Check for Visited States

How do we check if a state s has already been visited? Suppose we have a set Visited of visited states:

- 1 Simply check if $s \in \text{Visited}$.

Improving Check for Visited States

How do we check if a state s has already been visited? Suppose we have a set Visited of visited states:

- 1 Simply check if $s \in \text{Visited}$.
- 2 For all $t \in \text{Visited}$, check if $s \Leftarrow t$.

Improving Check for Visited States

How do we check if a state s has already been visited? Suppose we have a set Visited of visited states:

- 1 Simply check if $s \in \text{Visited}$.
- 2 For all $t \in \text{Visited}$, check if $s \Leftarrow t$.
 - Works, but **very expensive**.

Improving Check for Visited States

How do we check if a state s has already been visited? Suppose we have a set Visited of visited states:

- 1 Simply check if $s \in \text{Visited}$.
- 2 For all $t \in \text{Visited}$, check if $s \Leftarrow t$.
→ Works, but **very expensive**.

Can we do better than this?

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

- Start from an initial partition of W calculated wrt. labels.

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

- Start from an initial partition of W calculated wrt. labels.
- At each step, iterate through the elements, called **blocks**, B of the partition.

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

- Start from an initial partition of W calculated wrt. labels.
- At each step, iterate through the elements, called **blocks**, B of the partition.
- For $x \in B$ and $x' \in B'$, we say that B **sees** B' **via agent** i iff $xR_i x'$.

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

- Start from an initial partition of W calculated wrt. labels.
- At each step, iterate through the elements, called **blocks**, B of the partition.
- For $x \in B$ and $x' \in B'$, we say that B **sees** B' **via agent** i iff $xR_i x'$.
- For all agents $i \in Ag$, split B into sub-blocks B_1, \dots, B_k such that the worlds in each sub-block see the same blocks via agent i .

Ordered Partition Refinement

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement** for computing bisimulation contractions.

Standard partition refinement [PT87]. Let $s = ((W, R, L), w_d)$:

- Start from an initial partition of W calculated wrt. labels.
- At each step, iterate through the elements, called **blocks**, B of the partition.
- For $x \in B$ and $x' \in B'$, we say that B **sees** B' **via agent** i iff $xR_i x'$.
- For all agents $i \in \text{Ag}$, split B into sub-blocks B_1, \dots, B_k such that the worlds in each sub-block see the same blocks via agent i .
- Iterate until no more blocks can be split: the final partition is the **set of bisimulation equivalence classes of W** .

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .
- The **signature** a world $w \in W$ wrt. a partition (B_1, \dots, B_k) is defined as follows

$$\sigma_{(B_1, \dots, B_k)}(w) = L(w) \cup \{(i, n) \in Ag \times \mathbb{N} \mid \text{for some } v, wR_i v \text{ and } v \in B_n\}$$

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .
- The **signature** a world $w \in W$ wrt. a partition (B_1, \dots, B_k) is defined as follows

$$\sigma_{(B_1, \dots, B_k)}(w) = L(w) \cup \{(i, n) \in Ag \times \mathbb{N} \mid \text{for some } v, wR_i v \text{ and } v \in B_n\}$$

- Signatures give **unique identifiers** of worlds wrt. a partition.

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .
- The **signature** a world $w \in W$ wrt. a partition (B_1, \dots, B_k) is defined as follows

$$\sigma_{(B_1, \dots, B_k)}(w) = L(w) \cup \{(i, n) \in Ag \times \mathbb{N} \mid \text{for some } v, wR_i v \text{ and } v \in B_n\}$$

- Signatures give **unique identifiers** of worlds wrt. a partition.
- At each step, **blocks are split wrt. their signature**, until no more blocks can be split.

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .
- The **signature** a world $w \in W$ wrt. a partition (B_1, \dots, B_k) is defined as follows

$$\sigma_{(B_1, \dots, B_k)}(w) = L(w) \cup \{(i, n) \in Ag \times \mathbb{N} \mid \text{for some } v, wR_i v \text{ and } v \in B_n\}$$

- Signatures give **unique identifiers** of worlds wrt. a partition.
- At each step, **blocks are split wrt. their signature**, until no more blocks can be split.
- The **world-set** of the contraction of s is the **set of indices of the blocks in the final partition**.

Ordered Partition Refinement (Cont.)

Bolander et al. [BDH21] propose an improved algorithm called **ordered partition refinement (OPT)** for computing bisimulation contractions:

- Each block B_h is given a **numerical index** h .
- The **signature** a world $w \in W$ wrt. a partition (B_1, \dots, B_k) is defined as follows

$$\sigma_{(B_1, \dots, B_k)}(w) = L(w) \cup \{(i, n) \in Ag \times \mathbb{N} \mid \text{for some } v, wR_i v \text{ and } v \in B_n\}$$

- Signatures give **unique identifiers** of worlds wrt. a partition.
- At each step, **blocks are split wrt. their signature**, until no more blocks can be split.
- The **world-set** of the contraction of s is the **set of indices of the blocks in the final partition**.

Theorem ([BDH21])

If $s \sqsubseteq s'$, then the contractions computed by OPT are **identical**.

→ Bisimilarity check can be reduced to **identity check!**

Ordered Partition Refinement (Cont.)

Using ordered partition refinement, bisimilarity check can be reduced to **identity check!**

- An algorithm is provided to compute **policies** (mappings from states to actions) with a modified Pruning AND-OR (PrAO) search.
- Results show **improvements** both over a baseline planner that does not use OPT, and over the planner by Engesser et al. [Eng+17], a solver where each agent computes their policy distributively.

DEPTH-BOUNDED EPISTEMIC PLANNING

Depth-Bounded Epistemic Planning

In DEL-based epistemic planning agents can reason unboundedly about each other's knowledge.

- This leads to **undecidability** of the plan existence problem.
- Often unrealistic in many practical scenarios.

Depth-Bounded Epistemic Planning

In DEL-based epistemic planning agents can reason unboundedly about each other's knowledge.

- This leads to **undecidability** of the plan existence problem.
- Often unrealistic in many practical scenarios.

What if we **restricted the reasoning depth** of the planning agent to some bound b ?

- 💡 Reduce the size of epistemic states: **bounded bisimulation contractions**.
- 💡 Look for plans requiring the lowest bound: **iterative bound-deepening search**.

Bounded Bisimulations


💡 In a Nutshell: *b*-bisimilarity

- $x \Leftrightarrow_0 x'$ iff they agree on all propositional atoms.
- $x \Leftrightarrow_{b+1} x'$ iff $x \xrightarrow{i} y$ implies $x' \xrightarrow{i} y'$ and $x' \Leftrightarrow_b y'$ for some y' (and vice versa).

Proposition ([BRV01])

Two states are *b*-bisimilar iff they satisfy the same formulas up to modal depth *b*.

Example (Are x and x' bisimilar?)



Bounded Bisimulations


💡 In a Nutshell: b -bisimilarity

- $x \Leftrightarrow_0 x'$ iff they agree on all propositional atoms.
- $x \Leftrightarrow_{b+1} x'$ iff $x \xrightarrow{i} y$ implies $x' \xrightarrow{i} y'$ and $x' \Leftrightarrow_b y'$ for some y' (and vice versa).

Proposition ([BRV01])

Two states are **b -bisimilar** iff they **satisfy the same formulas up to modal depth b** .

Example (Are x and x' bisimilar? No, but they are 1-bisimilar!)




Rooted b -Contractions

Early definitions of bounded contractions in the literature did not behave as expected:

- **Standard b -contraction**: quotient structure of a model wrt. \cong_b .
- Standard b -contractions are in general **not minimal**.

Example (Standard (left) and minimal (right) b -contractions of a chain state)




- Each world of the chain can be **identified** by a formula of **modal depth** $\leq b$.
- Taking the quotient has no effect: we need to **keep all the worlds**.
- **Idea**: we only need to **keep some worlds**.

Rooted b -Contractions

Early definitions of bounded contractions in the literature did not behave as expected:

- **Standard b -contraction**: quotient structure of a model wrt. \cong_b .
- Standard b -contractions are in general **not minimal**.

Example (Standard (left) and minimal (right) b -contractions of a chain state)



- Each world of the chain can be **identified** by a formula of **modal depth** $\leq b$.
- Taking the quotient has no effect: we need to **keep all the worlds**.
- **Idea**: we only need to **keep some worlds**.



We improved the definition: **rooted b -contractions** guarantee **minimality**.

Canonical b -Contractions

Similar problem we had for standard contractions: **rooted b -contractions** of b -bisimilar states may be **non-isomorphic**!

→ Checking for visited states is **inefficient**.

Canonical b -Contractions

Similar problem we had for standard contractions: **rooted b -contractions** of b -bisimilar states may be **non-isomorphic**!

→ Checking for visited states is **inefficient**.



Improved definition called **canonical b -contractions**, based on the notion of **h -signatures**:

→ $\sigma_0(w) = (L(w), \emptyset)$

→ $\sigma_{h+1}(w) = (L(w), \Sigma_{h+1}(w))$, where $\Sigma_{h+1}(w)$ maps to each agent i a set

$$\Sigma_{h+1}(w, i) = \{\sigma_h(v) \mid wR_i v\}$$

→ Provide **unique identifiers** of h -bisimilar worlds.

Theorem (Identity [BBM25])

*The canonical b -contractions of b -bisimilar states are **identical**.*

From Breadth-First Search...

Let's start from a BFS with standard bisimulation contractions and check for visited states:

BFS

```
1: function BFS( $(s_0, Act, \phi_g)$ )
2:    $frontier \leftarrow \langle [s_0]_{\cong} \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow [s \otimes \alpha]_{\cong}$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

Proposition ([BRV01])

Two states are **bisimilar** iff they **satisfy the same formulas in $\mathcal{L}_{P, Ag}$** .

Proposition ([DHK07])

If $s \cong s'$ and α is applicable in both, then $s \otimes \alpha \cong s' \otimes \alpha$.

...To Bounded Search

Let b_0 be the **reasoning depth bound** of the planning agent (i.e., the agent can reason to formulas with **modal depth at most b_0**).

BoundedSearch

```
1: function BoundedSearch( $(s_0, Act, \phi_g), b_0$ )
2:    $frontier \leftarrow \langle [s_0]_{\leq} \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow [s \otimes \alpha]_{\leq}$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

...To Bounded Search

Let b_0 be the **reasoning depth bound** of the planning agent (i.e., the agent can reason to formulas with **modal depth at most b_0**).

BoundedSearch

```
1: function BoundedSearch( $(s_0, Act, \phi_g), b_0$ )
2:    $frontier \leftarrow \langle \llbracket s_0 \rrbracket_{b_0}^* \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow [s \otimes \alpha]_{\sqsubseteq}$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

Proposition ([BRV01])

Two states are **b -bisimilar** iff they **satisfy the same formulas up to modal depth b** .

...To Bounded Search

Let b_0 be the **reasoning depth bound** of the planning agent (i.e., the agent can reason to formulas with **modal depth at most b_0**).

BoundedSearch

```
1: function BoundedSearch( $(s_0, Act, \phi_g), b_0$ )
2:    $frontier \leftarrow \langle \llbracket s_0 \rrbracket_{b_0}^* \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow \llbracket s \otimes \alpha \rrbracket_{b_0}^* (?)$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

...To Bounded Search

Let b_0 be the **reasoning depth bound** of the planning agent (i.e., the agent can reason to formulas with **modal depth at most b_0**).

BoundedSearch

```
1: function BoundedSearch( $(s_0, Act, \phi_g), b_0$ )
2:    $frontier \leftarrow \langle \llbracket s_0 \rrbracket_{b_0}^* \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow \llbracket s \otimes \alpha \rrbracket_{b_0}^* (?)$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

Proposition ([BL22])

Let $s \Leftarrow_b s'$ and let α be an action with $md(\alpha) \leq b$. Then, $s \otimes \alpha \Leftarrow_{b-md(\alpha)} s' \otimes \alpha$.

Where $md(\alpha)$ denotes the **maximal modal depth** of all pre- and postconditions in α .

...To Bounded Search

Let b_0 be the **reasoning depth bound** of the planning agent (i.e., the agent can reason to formulas with **modal depth at most b_0**).

BoundedSearch

```
1: function BoundedSearch( $(s_0, Act, \phi_g), b_0$ )
2:    $frontier \leftarrow \langle \llbracket s_0 \rrbracket_{b_0}^* \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $s \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act$  applicable in  $s$  do
9:        $s' \leftarrow \llbracket s \otimes \alpha \rrbracket_{b_0}^*$ 
10:      If  $s'$  is not visited, push it to  $frontier$ 
11: return fail
```

Proposition ([BL22])

Let $s \Leftrightarrow_b s'$ and let α be an action with $md(\alpha) \leq b$. Then, $s \otimes \alpha \Leftrightarrow_{b-md(\alpha)} s' \otimes \alpha$.

Where $md(\alpha)$ denotes the **maximal modal depth** of all pre- and postconditions in α .

→ We need to **update** the bound value after an update.

Updating Bounds Value After Updates

We let a **node** of the search space be a pair $n = (s, b)$, where:

- 1 s is the **state** of n (denoted $n.state$).
- 2 b is the **(depth) bound** (denoted $n.bound$) → **maximum modal depth** of formulas we can safely evaluate in s .

Updating Bounds Value After Updates

We let a **node** of the search space be a pair $n = (s, b)$, where:

- 1 s is the **state** of n (denoted $n.\text{state}$).
- 2 b is the **(depth) bound** (denoted $n.\text{bound}$) → **maximum modal depth** of formulas we can safely evaluate in s .

In general, s will be a **b -contracted state** that can be thought of as an **approximation to the “real” state**.

→ We are always guaranteed that s is **at least b -bisimilar to the real state**.

Putting Everything Together

BoundedSearch

```
1: function BoundedSearch(( $s_0, Act, \phi_g$ ),  $b_0$ )
2:    $frontier \leftarrow \langle (\llbracket s_0 \rrbracket_{b_0}^*, b_0) \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $(s, b) \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act \mid b \geq md(\alpha)$  do
9:       if  $\alpha$  is applicable in  $s$  then
10:         $s' \leftarrow \llbracket s \otimes \alpha \rrbracket_{b-md(\alpha)}^*$ 
11:         $n' \leftarrow (s', b - md(\alpha))$ 
12:        If  $s'$  is not visited, push  $n'$  to  $frontier$ 
13: return fail
```

Proposition ([BRV01])

Two states are *b-bisimilar* iff they satisfy the same formulas up to modal depth b .

Proposition ([BL22])

Let $s \Leftarrow_b s'$ and let α be an action with $md(\alpha) \leq b$. Then, $s \otimes \alpha \Leftarrow_{b-md(\alpha)} s' \otimes \alpha$.

Putting Everything Together

BoundedSearch

```
1: function BoundedSearch(( $s_0, Act, \phi_g$ ),  $b_0$ )
2:    $frontier \leftarrow \langle (\|s_0\|_{b_0}^*, b_0) \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $(s, b) \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act \mid b \geq md(\alpha) + md(\phi_g)$  do
9:       if  $\alpha$  is applicable in  $s$  then
10:         $s' \leftarrow \|s \otimes \alpha\|_{b-md(\alpha)}^*$ 
11:         $n' \leftarrow (s', b - md(\alpha))$ 
12:        If  $s'$  is not visited, push  $n'$  to  $frontier$ 
13: return fail
```

Proposition ([BRV01])

Two states are *b-bisimilar* iff they satisfy the same formulas up to modal depth b .

Proposition ([BL22])

Let $s \Leftarrow_b s'$ and let α be an action with $md(\alpha) \leq b$. Then, $s \otimes \alpha \Leftarrow_{b-md(\alpha)} s' \otimes \alpha$.

Putting Everything Together

BoundedSearch

```
1: function BoundedSearch $((s_0, Act, \phi_g), b_0)$ 
2:    $frontier \leftarrow \langle \llbracket s_0 \rrbracket_{b_0}^*, b_0 \rangle$ 
3:    $visited \leftarrow \emptyset$ 
4:   while  $\neg frontier.empty()$  do
5:      $(s, b) \leftarrow frontier.pop()$ 
6:      $visited.push(s)$ 
7:     if  $s \models \phi_g$  then return plan to  $s$ 
8:     for all  $\alpha \in Act \mid b \geq md(\alpha) + md(\phi_g)$  do
9:       if  $\alpha$  is applicable in  $s$  then
10:         $s' \leftarrow \llbracket s \otimes \alpha \rrbracket_{b-md(\alpha)}^*$ 
11:         $n' \leftarrow (s', b - md(\alpha))$ 
12:        if  $s' \notin visited$  then  $frontier.push(n')$ 
13:   return fail
```

Proposition ([BRV01])

Two states are *b-bisimilar* iff they satisfy the same formulas up to modal depth b .

Proposition ([BL22])

Let $s \Leftarrow_b s'$ and let α be an action with $md(\alpha) \leq b$. Then, $s \otimes \alpha \Leftarrow_{b-md(\alpha)} s' \otimes \alpha$.

Theorem ([BBM25])

The canonical b -contractions of b -bisimilar states are *identical*.

Iterative Bound-Deepening Search

Iterative Bound-Deepening Search

```
1: function IBDS( $T = (s_0, Act, \phi_g)$ )
2:   for  $b \leftarrow md(\phi_g)$  to  $\infty$  do
3:      $\pi \leftarrow \text{BoundedSearch}(T, b)$ 
4:     if  $\pi \neq \text{fail}$  then return  $\pi$ 
```

We call **BoundedSearch** over increasing values of b :

- If $b < md(\phi_g)$, then the **bound is too low** to safely evaluate the goal formula.
- So initially we let $b = md(\phi_g)$.
- If no goal is found with bound b , we **increment the bound and try again**.

Improving Bounded Search

In a node $n = (s, b)$, the state s can be considered as an **approximation to modal depth b** of some “true state” t (namely, we are guaranteed that $s \trianglelefteq_b t$). However:

Improving Bounded Search

In a node $n = (s, b)$, the state s can be considered as an **approximation to modal depth b** of some “true state” t (namely, we are guaranteed that $s \trianglelefteq_b t$). However:

- In general it could be that $s \trianglelefteq t$!

Improving Bounded Search

In a node $n = (s, b)$, the state s can be considered as an **approximation to modal depth b** of some “true state” t (namely, we are guaranteed that $s \sqsubseteq_b t$). However:

- In general it could be that $s \sqsubseteq t$!
- In this case, when we update s with an action α , we don't have to decrease the bound.
 - Recall that **bisimilarity is preserved** after product update!

Improving Bounded Search

In a node $n = (s, b)$, the state s can be considered as an **approximation to modal depth b** of some “true state” t (namely, we are guaranteed that $s \sqsubseteq_b t$). However:

- In general it could be that $s \sqsubseteq t$!
- In this case, when we update s with an action α , we don’t have to decrease the bound.
 - Recall that **bisimilarity is preserved** after product update!

We can use this idea to include the following **optimizations** in BoundedSearch:

- We add a **third parameter** called *is_bisim* to our nodes, representing **whether the state of a node is bisimilar to its corresponding true state**.
- Depending on whether *is_bisim* holds, we **update a node with the appropriate bound value**.

Improving Bounded Search

In a node $n = (s, b)$, the state s can be considered as an **approximation to modal depth b** of some “true state” t (namely, we are guaranteed that $s \sqsubseteq_b t$). However:

- In general it could be that $s \sqsubseteq t$!
- In this case, when we update s with an action α , we don’t have to decrease the bound.
 - Recall that **bisimilarity is preserved** after product update!

We can use this idea to include the following **optimizations** in BoundedSearch:

- We add a **third parameter** called *is_bisim* to our nodes, representing **whether the state of a node is bisimilar to its corresponding true state**.
- Depending on whether *is_bisim* holds, we **update a node with the appropriate bound value**.
- Across different iterations of IBDS, we **preserve all nodes having *is_bisim* true**.
 - They would otherwise be **recomputed** in the next iteration!

Soundness, Completeness, Complexity

Let $T = (s_0, Act, \phi_g)$ be a planning task and let $b \geq md(\phi_g)$ be a constant.

Theorem (Soundness)

If **BoundedSearch**(T, b) returns an action sequence π , then π is a solution to T .


Theorem (Completeness)

If T has a solution of length ℓ , then **BoundedSearch**($T, c \cdot \ell + md(\phi_g)$) will find a solution to it, where $c = \max\{md(\alpha) \mid \alpha \in Act\}$.


Theorem (Complexity)

BoundedSearch runs in $(b+1)\text{-ExpTime}$.


IBDS vs. BFS




IBDS vs. Non-Optimized IBDS



Canonical vs. Rooted Contractions



IBDS vs. EFP 2.0



FUTURE DIRECTIONS

Many Ideas to Try Out

- DEL-based epistemic planning is a **hard problem**.
- Despite this, there have been **many recent promising advancements**.
- **Different ideas** have been explored, from compilation-based techniques, to heuristics, to bisimulation contractions.
- Many ideas haven't been tried yet!
 - **Symbolic** approaches, **SAT/SMT-based** epistemic planning, more heuristics.

One Language to Compare Them All

So many different frameworks, with many different semantics. How can we compare them?

- The **Epistemic Planning Domain Definition Language**.
- Combines a **PDDL-like syntax** with the **full DEL semantics**.
- **Different formalisms/fragments** can be define within the **same language!**
- Will soon be released!

Exciting News!



- See you all **next year** in Dublin for the **first Epistemic Planning Track** at the IPC!
- More details will be given at the **ICAPS Community Meeting** (**Thursday, November 13, at 15:30**).

Exciting News!



- See you all **next year** in Dublin for the **first Epistemic Planning Track** at the IPC!
- More details will be given at the **ICAPS Community Meeting** (**Thursday, November 13, at 15:30**).

Thank you! Questions?